

CLINICAL SUPPLY FORECASTING
DEMAND · SIMULATION · DISTRIBUTION

PREPARED FOR ACME BIOTECH
CONFIDENTIAL

ACME BIOTECH CLINICAL SUPPLY FORECASTING SIMULATOR

928801 v1.0 07 12 2026 PROPOSAL / ACME BIOTECH


A purpose-built demand forecasting and
scenario-planning engine that grows
with your team, and your clinical
programs.

PREPARED FOR

Acme Biotech

ISSUED

July 2026




Areas of opportunity

Where accuracy and speed
are being capped today.

Forecasting is innately high-leverage. Done well it compounds in your favor; done manually it compounds against you. In your current workflow, a few persistent, high-impact areas are quietly capping accuracy and speed:

- 01 **Forecasts live in spreadsheets that break as soon as enrollment, site count, or protocol assumptions change.**
- 02 **Demand planning is deterministic, so it cannot express the uncertainty that actually drives overage and stockout risk.**
- 03 **Scenario planning (enrollment swings, new sites, dosing changes) takes days of manual rework instead of minutes.**
- 04 **There is no single source of truth, so supply, clinical, and finance are each working from a different number.**




Clean pipelines, probabilistic modeling, repeatable SOPs.

02 / 10

Your solution

Closing these gaps is straightforward, and we have done it many times before. In practice the fix is almost always a combination of clean data pipelines, probabilistic modeling, and repeatable SOPs. Here is what that looks like for you:

- 01 **A parameterized forecasting model that recomputes instantly when any assumption changes.**
- 02 **Monte Carlo simulation that quantifies stockout and overage risk at your target service level.**
- 03 **One-click scenarios so you can pressure-test enrollment, site, and dosing changes in real time.**
- 04 **A shared, auditable dashboard that gives supply, clinical, and finance one aligned number, in demand and dollars.**



03 / 10

What ships

A working system,
not a slide deck.

Everything below is included in the engagement and itemized here
exactly as it will be delivered.

01	Custom demand forecasting model calibrated to your protocol & historical data.	INCLUDED
02	Monte Carlo risk engine with configurable service-level and buffer targets.	INCLUDED
03	Interactive scenario-planning dashboard (web-based, access-controlled).	INCLUDED
04	Data pipeline connecting your enrollment and inventory sources.	INCLUDED
05	Documented SOPs so your team can run and update forecasts without us.	INCLUDED
06	Handover training and 30 days of post-launch support.	INCLUDED

SCOPE MANIFEST · 6 LINE ITEMS

FIXED SCOPE

4 weeks, week by week

A phased build.
The timetable holds.

1

WEEK 1

Kickoff & Discovery

MAP THE TERRAIN BEFORE YOU BUILD

We read your current forecasting workflow end to end: existing spreadsheets, data sources, and the assumptions baked into each. Then we map exactly where the manual rework happens and where uncertainty is being dropped on the floor.

Every model input reflects how your program actually runs, not textbook assumptions, so the forecast is trustworthy from day one.

DATA SOURCE AUDIT · ASSUMPTION MAP · DAY-4 WORKING REVIEW

2

WEEK 2

Model Build

THE FORECASTING ENGINE, CALIBRATED TO YOU

We stand up the parameterized forecasting model and Monte Carlo risk engine directly against your historical data, not a generic template.

Calibrating against your real data means the risk numbers you see are the risk numbers you'll actually face, not generic benchmarks.

FORECASTING MODEL · MONTE CARLO RISK ENGINE · DAY-4 WORKING REVIEW

3

WEEK 3

Dashboard & Scenarios

ONE-CLICK ANSWERS FOR THE ROOM

We ship the interactive dashboard, wire up one-click scenario planning, and connect the whole thing to your enrollment and inventory sources.

A one-click scenario tool means you can defend supply decisions in the room, not after the meeting.

INTERACTIVE DASHBOARD · SCENARIO PLANNER · DATA PIPELINE CONNECTED

4

WEEK 4

Signal & Warm Handoff

PROOF IT HOLDS, THEN A CLEAN HANDOFF

We run the model against known outcomes so you see real signal that it holds, not performance we claim on our own docs. Then we document the SOPs your team needs and walk them through owning it.

Seeing the model checked against known outcomes before handoff means you inherit a proven system, not an unproven black box.

SIGNAL REPORT · DOCUMENTED SOPS · TEAM TRAINING · 30-DAY SUPPORT

The engagement

Senior team.
Narrow focus.

Two operators, partnered with your team, 4 weeks.

Two operators work directly with your team for 4 weeks: one leads the engagement end to end; the other builds the forecasting model and integrates it against your live data. We run this with you, not for you: one question, 4 weeks, and it is the only thing we work on.

CORE

Engagement Lead

On the work full time from kickoff to handoff. Primary point of contact for the engagement.

CORE


Forecasting & Data Systems

Builds the model and Monte Carlo risk engine, and integrates it against your live data sources.

SCOPE

Fixed scope

4 build phases over 4 weeks, described in full in the plan.



Why OPFOR

Modeling depth,
operational SOPs.

6+ yrs

APPLYING AI TO SUPPLY AND
FORECASTING

95%+

TARGET SERVICE LEVELS
ENGINEERED INTO EVERY MODEL

End-to-end

FROM DATA PIPELINE TO
BOARDROOM DASHBOARD

We build AI and forecasting systems for supply-constrained operations. Our work has driven measurable reductions in overage and stockout risk across clinical and commercial supply programs. We pair modeling depth with practical SOPs, so the system keeps working after we leave.

01 **An assumption change becomes an answer in seconds.**

Any assumption change re-forecasts in seconds, in your browser. Self-serve, unlimited, no analyst queue.

03 **Calibrated to your data, not a benchmark.**

Every model ships trained against your actual historical and protocol data, not a generic industry template.

05 **Decision-ready, or we say so.**

If the data can't support a confident forecast at your target service level, we tell you before you pay for a model that can't deliver it.

02 **IRT-agnostic by architecture.**

A neutral layer over whatever randomization system each study runs. Medidata live today; other formats onboard per study.

04 **Every forecast is reproducible.**

Deterministic engine: any forecast re-derives exactly from its configuration snapshot. Every change is versioned.

06 **No patient data. Ever.**

No subject-level records, no treatment arms. Enrollment reduces to per-site counts in your browser. Only the counts leave your machine.

DEMAND PER	PROTOCOL	PER ARM	PER VISIT SCHEDULE	ENROLLMENT CURVES	TITRATION PATHS
DROPOUT	MODELING	DEPOT ALLOCATION	SHELF-LIFE AWARE	LIVE REFORECAST	IRT ACTUALS
COHORT	MANAGEMENT	MONTE CARLO	SIMULATIONS	DEMAND AGGREGATION	BUDGET PLANNING
					DP ALIGNMENT

Data boundaries


What we hold.
What we never touch.

OPFOR is planning and decision support. The boundary below is architectural, not policy.

- | | | |
|----|---|---|
| 01 | Zero access to your systems | Zero OPFOR accounts, credentials, or permissions exist in your IRT or any validated system. Your team exports the reports it already receives and uploads them. That is the entire data path. |
| 02 | Your protocol never enters the product | There is no document upload. Setup consumes planning parameters only: countries, visit schedule, kit types. Drafting happens with you, outside the product. Blinded data is firewalled the same way: reduced in your browser, on your machine, and never transmitted to us or stored. |
| 03 | Counts, not people | The subject report is reduced to per-site, per-month counts in your browser. Only the counts leave your machine. No individual subject record is stored. |
| 04 | No treatment-arm data, in any mode | The parser never reads arm, stratum, or randomization columns. The forecast runs on kit-type demand, so blinding stays intact by construction, not by policy. |

Out of GxP scope by design

OPFOR is never a system of record and performs no GxP activity. It ingests only decision-support exports your team already receives. No source documents. Adoption does not wait on a computerized-system validation project.



08 / 10

The investment

Seconds to an answer.
No vendor queue.

Clinical Supply Forecasting Simulator	\$45,000
White-glove onboarding	\$4,500
QTY 3	
Training hours	\$2,000
QTY 5	
IRT integration OPTIONAL	\$5,680 /EA
Optional. Billed per integration; refunded onto your following-year renewal.	

TOTAL INVESTMENT

\$51,500

The bonuses, yours with the build.

Locked-for-life price guarantee: your annual investment never increases	\$3,250	GUARANTEED
Testimonial credit off your next renewal	\$7,000	ON RENEWAL
Unlimited scenario breakdown videos	\$2,950	INCLUDED
Weekly working call, one hour a week	\$14,040	INCLUDED
SOC 2-compliant host-for-you	\$749	INCLUDED
(2) RTSM/IRT integration fee waivers (\$5,680 each)	\$11,360	ON RENEWAL
Unlimited Anthropic API usage	\$1,200	INCLUDED
Step-by-step study build master manual	\$790	INCLUDED
Long Range Planning module	\$4,795	INCLUDED
Drug Product DOM Alignment module	\$3,796	INCLUDED
24/7 customer service access		INCLUDED
(21) Master-builder training video guides	\$5,229	INCLUDED
(11) Do-it-for-you white glove study builds	\$8,217	INCLUDED

TOTAL BONUS VALUE

\$63,376

FREE

50% AT SIGNATURE · BALANCE ACROSS Q3 AND Q4 · QUARTERLY THEREAFTER · NET 30

09 / 10

Next steps

Three moves.
Then we build.

1

Approve this proposal to reserve
your build slot.

2

Kickoff call to grant data access
and confirm assumptions.

3

We begin the Discovery &
Scenario Build within one week of
kickoff.


THE GUARANTEE

If the delivered model does not hold to the agreed service-level target on
validation data, we keep working at no additional cost until it does.

READY WHEN YOU ARE

SEE IT RUN. SUPPLY THE WORLD.

INFO@OPFORSUPPLY.COM · OPFORSUPPLY.COM



Services agreement

Executed off-platform
by client legal.

This Agreement is entered into as of July 12, 2026 (the "Effective Date") by and between OPFOR ("Provider") and Acme Biotech ("Client"), together the "Parties."

- 1 **Services:** Provider will perform the services described in this proposal (the "Services") in accordance with the scope, timeline, and deliverables set out above.
- 2.1 **Fees:** Client will pay the fees set out in the Investment section above.
- 2.2 **Invoicing:** Provider will invoice fifty percent (50%) of first-year fees upon execution of this Agreement, and the remainder in two equal installments invoiced in the third and fourth calendar quarters of the first year. Thereafter, fees are invoiced quarterly in advance. Payment is due within thirty (30) days of the invoice date.
- 3.1 **Payment Terms:** Late payments accrue interest at 1.5% per month, or the maximum allowed by law, whichever is lower.
- 3.2 **Taxes:** Client is responsible for all sales, use, VAT, or similar taxes, excluding taxes on Provider's net income.
- 3.3 **Third-Party Costs:** Client will reimburse Provider for third-party software licenses, API-usage fees, or comparable platform costs necessary to perform the Services, provided Provider obtains Client's written approval before incurring any such expense.
- 4 **Client Responsibilities**
 - Timely access to relevant staff, systems, data, and credentials.
 - A single point of contact with decision-making authority.
 - Prompt review/approval of deliverables (deemed accepted if no written rejection within five (5) business days).
 - Compliance with all laws governing Client data supplied to Provider.
- 5.1 **Provider Materials:** Tools, code libraries, frameworks, models, prompts, pre-existing IP, and generic know-how ("Background IP") remain Provider property; Provider grants Client a non-exclusive, perpetual, worldwide license to use Background IP only as embedded in Deliverables.
- 5.2 **Deliverables:** Custom code, configurations, documentation, and model outputs created specifically for Client under this proposal ("Deliverables") become Client property upon full payment.
- 6.1 **Confidentiality:** Each Party will protect the other's Confidential Information with at least the same care it uses for its own similar information, but no less than reasonable care.

- 6.2
- Survival: Obligations last three (3) years after termination, except trade secrets, which remain protected while they are trade secrets under applicable law.
- 7.1
- Provider Indemnity: Provider will defend and indemnify Client against third-party claims that Deliverables infringe intellectual-property rights.
- 7.2
- Client Indemnity: Client will defend and indemnify Provider against claims arising from Client data, misuse of Deliverables, or Client's violation of law.
- 8
- Limitation of Liability: Each Party's total liability is limited to the greater of (i) US \$100,000 or (ii) the fees paid or payable under this Agreement in the twelve (12) months preceding the claim. Neither Party is liable for indirect, incidental, consequential, special, or punitive damages, including lost profits or revenue.
- 9
- General: Provider is an independent contractor. This Agreement is governed by the laws of [Governing State/Country], without regard to conflict-of-laws principles. This Agreement is the entire understanding between the Parties and supersedes all prior proposals or agreements. Amendments must be in writing signed by both Parties.

IN WITNESS WHEREOF, the Parties have executed this Agreement as of the Effective Date.

CLIENT

Acme Biotech

NAME / SIGNATURE

DATE

SERVICE PROVIDER

OPFOR

NAME / SIGNATURE

DATE

OPFOR SUPPLY · CLINICAL SUPPLY FORECASTING

10 / 10

+ BSL